

The City Understands the Place:

Location Memory from Infrastructure Cameras as a Prior for Autonomous Driving

Gamzat Gadzhibutaev*

Working paper — draft of 29 September 2026

Abstract

A robotaxi perceives a place only while it is there; a fixed camera watches the same place for hours, days and months. We ask whether this long-term *location memory* is worth anything to an autonomous vehicle, and when it becomes a liability. Memory is represented as exposure-normalised sufficient statistics with uncertainty—where pedestrians step onto the road, how traffic turns, where vehicles brake hard—and is evaluated in three steps. (i) On two hours of real trajectories from 12 fixed 4K infrastructure cameras in Washington, DC (TGSIM Foggy Bottom), memory built from the first 79 minutes predicts the last 40: 73 % of held-out pedestrian road entries fall inside the 10 % of road area it ranks highest (2.27 bits per entry over a uniform prior), and for vehicles that change direction it picks their heading 3 s ahead correctly 59 % of the time, against 19 % for area-wide statistics. (ii) In a closed-loop occlusion simulation calibrated on the same data, an occlusion-aware planner with 10 hours of memory has 26 % fewer pedestrian collisions at equal trip time (95 % CI 20–33 %) than a planner that knows only the street’s average pedestrian rate, matching perfect knowledge; the benefit grows with the spatial unevenness of pedestrian activity, from 0 to 56 %. Two failure modes are severe: stale memory is worse than none, and a live camera that fails silently is the most dangerous state studied. Letting memory only *add* caution, and expiring live information with a heartbeat, remove most of the harm. (iii) Treating every real vehicle in the data as a potential fleet car shows that a fleet can learn the same memory, but at the pace of equipped cars passing by: at 10 % penetration it needs roughly four times the calendar time of a fixed camera.

1 Introduction

Automated vehicles perceive the world with their own sensors, here and now. At an urban kerb lined with parked cars, this is a hard limit: a pedestrian stepping out between two vans is invisible until the last moment, and a planner that must stop for every pedestrian who *might* be there crawls along the block. Cities, meanwhile, already operate thousands of fixed cameras, and the number keeps growing. Vehicle-to-infrastructure (V2I) cooperative perception uses such sensors to share *live* detections with vehicles [1–4], and standards for doing so exist [5, 6].

A fixed sensor also has something a passing vehicle cannot have: *temporal depth*. It watches the same place continuously, and over hours and days learns where people cross, how traffic turns, where drivers brake. A robotaxi may arrive at an intersection for the first time, yet the city may already hold the equivalent of thousands of past passages through it. In short: the vehicle understands its surroundings; the city understands the place.

This paper asks whether such location memory is worth anything to an autonomous vehicle. Following a long-standing argument that part of vehicle intelligence can live in the infrastructure [7], we treat the city’s knowledge strictly as a *probabilistic prior* for the vehicle’s planner, never as

*HSE University, Faculty of Economics, Nizhny Novgorod, Russia. gbgadzhibutaev@edu.hse.ru

a command: the vehicle’s onboard perception keeps the final decision. The questions are therefore empirical: does memory predict what happens later, does it change driving outcomes, when does it fail, and why would fixed cameras be needed if a fleet of vehicles could learn the same thing?

Contributions.

1. A compact, identity-free representation of location memory as exposure-normalised sufficient statistics with empirical-Bayes uncertainty, and a message design whose cost is kilobytes rather than video (§3).
2. Evidence on real infrastructure-camera data that memory predicts held-out behaviour, together with its honest limits: rare conflicts are statistically invisible in two hours, and signal-driven stop-and-go dominates vehicle prediction error (§4).
3. A closed-loop, rare-event simulation calibrated on the same data that quantifies the safety–time benefit, its dependence on the unevenness of pedestrian activity and on camera coverage, and two failure modes with mitigations (§5).
4. A fleet-versus-infrastructure comparison on real traffic showing that fleet memory is limited by exposure rather than by quality (§6).

All code, tests and results are in the accompanying repository and reproduce on a laptop in about twenty minutes.

2 Related work

Cooperative perception and V2X. Vehicle–infrastructure datasets [1, 2] and methods [3, 4] fuse live roadside and onboard observations; V2X-Seq shows that infrastructure trajectories improve forecasting, but with seconds, not months, of history. ETSI’s Collective Perception Service [5] standardises the exchange of detected objects, and the Local Dynamic Map [6] already layers static, quasi-static, transient and highly dynamic information—the layered “city model” of this paper is an extension of that structure, not a new architecture. Large deployments are under way, e.g. China’s vehicle–road–cloud pilots in 20 cities [8]; a recent survey organises the field [9]. Closest in spirit is a 2026 vision paper arguing that fixed roadside sensors offer “temporal depth” and proposing infrastructure-centric world models [10]; it contains no experiments.

Memory from vehicle fleets. Crowd-sourced maps store aggregate driving behaviour per road segment [11]; CHAMP maps pedestrian zones from repeated passes of instrumented vehicles [12]; past traversals [13] and neural map priors [14] improve perception. These are the natural competitors of infrastructure memory, and §6 compares the two on the same traffic.

Maps of dynamics. Mobile robotics represents place-dependent motion patterns as maps of dynamics [15]; CLiFF-map-based prediction [16] and its time-conditioned extension [17] predict long-term human motion. Our flow layer is a discrete CLiFF map.

Occlusion-aware planning. Planners place phantom road users in occluded areas, which is safe but conservative; V2X information reduces the uncertainty and the conservatism [18]. Emergence probabilities of hidden pedestrians have been estimated from the current context [19]; we estimate them from long-term infrastructure history.

Safety analytics from infrastructure. Commercial systems accumulate near-miss statistics from existing cameras [20], roadside LiDAR supports post-hoc safety auditing [21], and long-term roadside observations have been used to learn realistic, location-specific driving environments for AV testing [22]. In these works memory serves agencies or test engineers; here it serves the vehicle’s planner.




Figure 1: Architecture. Video never leaves the edge; the city model keeps a live state, a long-term location memory and short-term predictions, and sends each vehicle only the part relevant to its corridor, with an age and a confidence. The vehicle treats it as a prior.

3 Location memory

Observations. Fixed sensors observe tracks $o = (\text{id}, t, \mathbf{x}, \mathbf{v}, \text{class})$ in a common map frame. Memory is kept per grid cell c (1 m) and per lane or crosswalk segment r , and consists of *sufficient statistics*, never raw trajectories or identities: long-term mobility traces are re-identifiable even when anonymised [23], while per-cell counts are not.

Layers. For each road-user class (pedestrian, micro-mobility, vehicle) the memory stores (a) *occupancy and exposure*: agent-seconds and distinct passages per cell, used to normalise every rate; (b) a *flow field*: for 8 heading bins per cell, the count, mean direction and mean speed of moving samples (a discrete CLiFF map [16]); (c) *manoeuvre transitions*: counts $T_c[b_0, b_3]$ of heading bin now versus 3 s later; (d) vehicle *speed quantiles* per cell; and events: (e) *pedestrian emergence*, the first appearance of a pedestrian on the road after stitching tracker fragments; (f) *hard braking*, sustained deceleration $\leq -3.5 \text{ m/s}^2$ for $\geq 0.5 \text{ s}$ at $\geq 3 \text{ m/s}$; (g) vehicle–pedestrian *conflicts*, time-to-collision [24] $\leq 1.5 \text{ s}$ under constant velocity with a three-disc vehicle footprint.

Rates with uncertainty. For segment r with k_r events in n_r passages, the per-passage probability has an empirical-Bayes posterior

$$p_r \sim \text{Beta}(m\bar{p} + k_r, m(1 - \bar{p}) + n_r - k_r), \quad \bar{p} = \sum_r k_r / \sum_r n_r, \quad (1)$$

with prior strength m chosen by beta-binomial marginal likelihood, which shrinks sparsely observed segments towards the pooled rate. Manoeuvre distributions are shrunk the same way towards area-wide transitions, $\hat{p}(b_3 | c, b_0) = (T_c[b_0, b_3] + \alpha \bar{p}(b_3 | b_0)) / (\sum_b T_c[b_0, b] + \alpha)$ with $\alpha = 2$, and the emergence intensity $\lambda(\mathbf{x})$ is a Gaussian-smoothed ($\sigma = 1.5 \text{ m}$) count per unit area and time over the observed road area.

The message. Layers are split by how often they change. The memory tile for a vehicle’s corridor (90 m long, 24 m wide, six one-byte layers) is sent once and cached; only the live object list and short-term predictions stream at 10 Hz. A seemingly compact alternative—a dense $200 \times 200 \text{ m}$ tensor at 0.5 m resolution with ten channels at 10 Hz—costs about 128 Mbit/s, as much as raw video; “not video” is not automatically cheap.

4 Part 1: does memory predict the place?

Data. TGSIM Foggy Bottom [25] was recorded by 12 stationary 4K infrastructure cameras covering four intersections, crosswalks and the road segments between them in Washington, DC, over two hours of a weekday afternoon at 10 Hz, with all road-user classes and one SAE Level 3 test vehicle. After stitching pedestrian fragments (a track starting $\leq 1 \text{ s}$ after another ended within 2 m of its constant-velocity extrapolation; $15,307 \rightarrow 14,237$ tracks) and removing tracks shorter than 1 s, 2.48 M points remain: 11,521 pedestrian and 4,930 vehicle tracks, with 11,510 pedestrian emergences, 1,424 hard-braking onsets and 36 conflicts.




Figure 2: Location memory built from the first 79 minutes: (a) vehicle flow and typical speed, (b) where pedestrians step onto the road, (c) posterior hard-braking probability per lane segment, (d) the 36 vehicle–pedestrian conflicts of the two hours.


Protocol. Memory is built from the first 79 minutes and evaluated on the last 40, which it never saw; blocked three-fold cross-validation holds out each 40-minute block once. The baseline is always what a vehicle arriving for the first time would know: a uniform distribution over the road area, or area-wide pooled statistics. Scores are information gain in bits, capture curves, Brier skill, ROC AUC and top-1 accuracy, with bootstrap confidence intervals; prediction hyper-parameters are tuned on a separate validation split.

Where people step onto the road. Figure 2 shows the memory; Table 1 and Figure 3 summarise the held-out scores. Seventy-three percent of held-out pedestrian entries fall inside the 10 % of road area that memory ranks highest, 2.27 bits per entry better than a uniform prior; cross-validation folds agree (2.27–2.44 bits). The harder case—entries outside crosswalks, between parked or queued cars—is still captured four times better than uniform. The learning curve is not saturated: 5, 10, 20, 40 and 79 minutes of history give 1.46, 1.73, 1.93, 2.15 and 2.27 bits.

How traffic behaves there. Segment-level hard-braking rates are calibrated and discriminative (AUC 0.85; cross-validated Brier skill +0.085 to +0.115). Memory also knows the manoeuvre distribution of each place: for vehicles whose heading bin changes within 3s it predicts the new heading correctly 59 % of the time, versus 19 % with area-wide transitions (Figure 4a); for pedestrians the gain is smaller (0.10 bits per sample). Used deterministically, the flow prior reduces the 3s error of turning vehicles by 18 % but cannot improve straight-going ones: a single rollout must commit to one manoeuvre, so the planner should receive the distribution rather than one guess (Figure 4b).

Table 1: Part 1: memory versus the location-agnostic prior on the held-out 40 minutes.

Question (held-out sample)	Memory	No memory
Where do pedestrians step onto the road? (4,334 entries) ...outside crosswalks (799)	73 % in the top 10 % of area; +2.27 bits [2.22, 2.31] 42 % in the top 10 %; +0.84 bits [0.72, 0.97]	10 % (uniform) 10 %
Which lane segments produce hard braking? (9,905 passages)	AUC 0.85; Brier skill +0.08	pooled rate 5.7 %
Heading in 3 s of vehicles that change direction (17,154)	59 % correct; $P(\text{true}) = 0.52$	19 %; 0.33
3 s position error of turning vehicles (1,110)	5.22 m (−18 % [−16, −20])	6.38 m (constant velocity)
Conflict per pedestrian passage (5,601; 8 events)	Brier skill −0.005 (no gain)	pooled rate 0.23 %

**Figure 3:** (a) Share of held-out pedestrian road entries captured by the top-ranked share of road area. (b) Information gain versus minutes of history.

What memory does not know. Two negatives are informative. First, vehicle–pedestrian conflicts are rare—36 in two hours, eight in the test window—and a per-crosswalk conflict prior is indistinguishable from the pooled rate. Rare safety events need weeks to months of memory, which is an argument for long-term infrastructure memory, not against it. Second, 3.2 of the 3.3 m constant-velocity error of vehicles at 3 s is longitudinal: stop-and-go at signals. A memory of the place cannot know the current signal phase, which only live data can supply.

Occlusion and bandwidth. Treating every moving vehicle as a virtual robotaxi with a roof sensor and 2-D line of sight blocked by other vehicles (with stopped vehicles keeping their orientation), 10.4 % of 3,042 encounters with a pedestrian on or entering the vehicle’s path start hidden, and in 3.8 % the infrastructure knows about the pedestrian at least 1 s earlier (lower bound with buses and trucks as the only occluders: 0.7 % and 0.3 %). These numbers are strongly underestimated because the data contain neither sidewalks nor buildings. For the real test vehicle at one such moment, the corridor tile is 2.7 KB and the live object list about 12 kbit/s, against roughly 180 Mbit/s of raw 4K video from the 12 cameras.




Figure 4: (a) Heading 3 s ahead for road users that change direction. (b) Final displacement error of a single 3 s rollout steered by the flow memory, tuned on validation data for curved motion.

5 Part 2: does memory help a robotaxi drive?

5.1 Simulation calibrated on the real block

Scene. A 90 m block with a dense parked row (cars 4.3–5.2 m, gaps mostly 0.6–2 m); the robotaxi drives in the adjacent lane with 1 m lateral clearance at up to 25 mph (11.2 m/s). Pedestrians step out between parked cars *where and how often they did in TGSIM*: the emergence profile $\lambda(x)$ is the smoothed histogram of the 99 real mid-block emergences on the block between two of the intersections (about 50 per hour, peak 3.3 times the mean), and walking speeds follow the real distribution of those pedestrians (median 1.49 m/s, 90th percentile 2.25 m/s). Ninety percent of pedestrians are attentive (they wait if the car is less than 3 s away), ten percent step out without looking, and nobody walks into the side of a passing car.

Vehicle and perception. A roof sensor sees a pedestrian when its 2-D sight line clears the parked row; the planner reacts after $t_r = 0.25$ s of continuous visibility and can brake at $a_e = 7$ m/s². An optional infrastructure camera reports pedestrians approaching the kerb with 0.3 s latency.

Planner. A pedestrian leaving the occlusion edge at distance d ahead of the bumper is unavoidable if the car can neither stop before it nor pass before the pedestrian reaches its path: $vT_p < d < vt_r + v^2/(2a_e)$, where T_p is the pedestrian’s time from the occlusion edge to the vehicle’s swath. With emergences at rate $\lambda(x)$ per metre and second, the car is exposed at x for

$$b(v) = \max(0, t_r + v/(2a_e) - T_p) \quad (2)$$

seconds, and minimising $\int dx/v + w \int \lambda b(v) dx$ pointwise yields the target speed

$$v^*(x) = \min\left(v_{\max}, \max(v_{\text{wc}}, \sqrt{2a_e/(w\lambda(x))})\right), \quad v_{\text{wc}} = 2a_e(T_p - t_r), \quad (3)$$

smoothed backwards for comfortable deceleration; v_{wc} (9 km/h here) is the worst-case speed at which $b = 0$. All planners share this model and the exchange rate w between time and risk; they differ *only* in the intensity map λ they believe (Table 2). The no-memory planner knows the street’s true *average* rate, so any gain measures the value of knowing *where*, not *how many*.

Rare-event design. Each episode contains exactly one pedestrian, who reaches the kerb within a window of $W = 11$ s around the robotaxi’s nominal arrival at that spot. Per-traversal rates are

$$\mathbb{E}[\text{events per traversal}] = \Lambda W [p_i P(\text{event} \mid \text{inattentive}) + (1 - p_i) P(\text{event} \mid \text{attentive})], \quad (4)$$

Table 2: Planners in the simulator. All use Eq. (3) with the same w .

Planner	Believed pedestrian intensity
Worst case	a pedestrian may step out anywhere at any time ($v = v_{wc}$ along the parked row)
No memory	the street’s true mean rate, uniform along the block
Location memory	estimated from H hours of simulated camera observation (10 h: 448 events, with sampling noise)
Memory, adds caution only	memory, but never faster than the no-memory planner
Perfect memory	the true $\lambda(x)$ (upper bound)
+ live camera	covered sections believed clear; approaching pedestrians reported




Figure 5: (a) Collisions versus trip time; each dot is one caution level. The arrows compare planners at the circled no-memory operating point. (b) Near misses (braking $\geq 4 \text{ m/s}^2$ needed). (c) Real emergence profile, memory estimate and resulting target speeds.

with Λ the street’s emergence rate and $p_i = 0.1$. Every configuration runs 30,000 inattentive and 3,000 attentive episodes, identical across planners (common random numbers); differences carry paired-bootstrap intervals. A unit test checks that the worst-case planner never hits a pedestrian walking at the planner’s design speed, i.e. that planner model and simulator agree. Because the inattentive share is a free parameter, *absolute* collision rates are not calibrated to crash statistics; only comparisons between planners are meaningful.

5.2 Results

Memory moves the safety–time frontier. At the trip time of the no-memory planner at a human-like caution level, the planner with 10 hours of memory has 26.5 % fewer collisions (95 % CI 20.1–32.8 %), equal to perfect memory (26.7 %); at equal risk it is 1.1 s (6 %) faster over the block, and near misses also drop (Figure 5). It slows to 17 km/h at the hotspot and drives 40 km/h where nobody steps out, instead of a uniform 29 km/h; the worst-case planner crawls at 9 km/h. The benefit saturates at about 4 hours of history (≈ 200 observed emergences); 15 minutes of history is sometimes worse than no memory (Figure 6b).

It pays off where activity is uneven. Under a risk proportional to λv and a time proportional to $1/v$, the optimal speed scales as $\lambda^{-1/2}$ and the achievable risk–time product relative to ignoring location is $(\mathbb{E}\sqrt{\lambda})^2/\mathbb{E}\lambda$. We use $1 - (\mathbb{E}\sqrt{\lambda})^2/\mathbb{E}\lambda$ as a heterogeneity index of a place. Redistributing the same total pedestrian rate along the block, the collision reduction at equal time grows monotonically with this index: 0 % for a uniform street, 27 % for the real block (index 18 %), 56 % when the peak is ten times the mean (Figure 6c). The index is a predictor rather than a




Figure 6: (a) Collisions at equal trip time versus the share of the block seen by cameras. (b) Versus hours of history. (c) Collision reduction versus the heterogeneity of pedestrian activity; bands are 95 % paired-bootstrap intervals.




Figure 7: Collisions per 10,000 traversals at the same caution setting. (a) The hotspot moved and memory did not. (b) Camera outage with and without a heartbeat.

bound—the real planner’s risk has an offset v_{wc} —but it is measurable from a few hours of camera data and tells a city where memory is worth deploying.

Live cameras are the bigger lever. With a camera covering the whole block, collisions fall by about 99 % and the block is driven 3s faster; memory then adds nothing. With partial coverage memory still reduces collisions by 10–36 %, depending on how uneven the *uncovered* part is (Figure 6a).

Failure modes. Two failures are severe (Figure 7). When the hotspot moves 21 m to the quietest part of the block and memory does not know, memory trusted in both directions drives fast exactly where people now step out: 2.52 collisions per 10,000 traversals versus 2.00 without memory. A caution floor barely helps (2.44). If memory may only *add* caution, the stale memory still beats no memory (1.62, at the cost of 1.4s); with correct memory this restriction keeps most of the benefit (–21 % instead of –26 % at equal time). When the camera fails and the vehicle is not told, it keeps trusting an empty street: 3.83 collisions per 10,000, the most dangerous state in the study and worse than no infrastructure at all. A heartbeat that reverts the vehicle to its memory prior brings this to 1.54. Finally, with the camera up the simple planner drives near the limit and brakes hard twice as often (101 versus 50 per 1,000 traversals), because attentive pedestrians accept 3s gaps: the comfort side of live data needs its own planning.

Table 3: Fleet versus infrastructure memory on the same real traffic. Time to full benefit: simulator, two memory draws per setting (indicative).

Memory source (equipped vehicles / h)	Pedestrians seen	Gain: all / off-crosswalk	Time to full benefit
Fixed cameras	100 %	2.27 / 0.84 bits	≈ 1 h
Fleet, 100 % (184/h)	98 %	2.24 / 0.81 bits	≈ 1 h
Fleet, 30 % (≈ 55 /h)	71 %	2.14 / 0.71 bits	—
Fleet, 10 % (≈ 18 /h)	37 %	1.96 / 0.62 bits	≈ 4 h
Fleet, 3 % (≈ 5 /h)	17 %	1.62 / 0.37 bits	≈ 16 h




Figure 8: Fleet memory on real traffic: (a) share of pedestrians seen, (b) held-out information gain versus penetration, (c) learning speed of camera and fleet memory.

6 Part 3: why fixed cameras, if a fleet can remember?

A fleet memory only knows what passing vehicles happened to see, and records a pedestrian where and when a vehicle first saw it. We treat each of the 4,828 driving vehicles in TGSIM as a potential fleet car with a 50 m roof sensor whose 2-D line of sight can be blocked by other vehicles; a random share of them (the penetration) is equipped. The fleet memory is built from equipped vehicles’ first sightings in the first 79 minutes and scored on the same held-out window as the camera memory. In the simulator, the share of time each metre of kerb is watched by equipped vehicles—measured on the real traffic of the block, 184 driving vehicles per hour—thins the observations from which memory is estimated.

When a car is present it sees the pedestrian almost as well as a camera: median 0.2 s late and 0.2 m from the true entry point (Table 3, Figure 8). What limits a fleet is exposure—equipped vehicles per hour times hours. A 10 % fleet’s memory after 79 minutes (1.96 bits) matches the camera memory after roughly 30 minutes, and in the simulator it needs about 4 hours to deliver the full planning benefit that a camera delivers in about one; a 3 % fleet needs about 16 hours and, after one hour, is worse than no memory. Entries outside crosswalks, where hidden-pedestrian risk lives, are where the fleet lags most. Correcting for per-location exposure made no difference here because equipped vehicles watched the kerb fairly uniformly (19–36 % of the time at 10 %), but it matters where coverage is uneven. This rush-hour block is the fleet’s best case; the fleet is idealised (perfect detection within 50 m, every sighting uploaded) and can only observe pedestrians the cameras also tracked, so the fleet numbers are upper bounds.

7 Discussion

Design rules. Four rules follow from the experiments. (1) *A prior, not an oracle*: the vehicle decides; the city informs, with an age and a confidence on every message. (2) *Asymmetric trust*: a location prior may add caution freely, while removing caution requires fresh evidence; this costs about a fifth of the benefit when memory is right and prevents harm when it is wrong. (3) *Trust*

must expire: live “the area is clear” information is the most valuable and the most dangerous message; a missing heartbeat must revert the vehicle to its prior. (4) *Invest where places are uneven*: the benefit of memory tracks a heterogeneity index measurable from hours of camera data.

Cameras and fleets are complementary. Fixed cameras give memory in hours instead of days, independent of traffic volume, provide live perception beyond a vehicle’s line of sight, and accumulate the exposure needed for rare events. Fleets cover places without cameras and can keep a camera’s memory fresh; a hybrid memory with change detection is the natural next system.

Limitations. Part 1 uses two hours of one sunny weekday afternoon at one site; weather, lighting, day-of-week and cross-day transfer cannot yet be tested, and rare events are out of reach. The data contain road users only inside road polygons, and neither sidewalks nor buildings, so occlusion and infrastructure lead times are underestimated. The simulator is stylised: a one-dimensional longitudinal planner, 2-D line of sight, one pedestrian per episode, a synthetic parked row, pedestrians from the near side only and a free inattentive share; the camera never raises false alarms, which flatters live infrastructure. Memory learned from human behaviour is a prior, not a policy: that most drivers do something at a place does not make it safe.

Next steps. Repeated recordings of the same intersections on different days [26] and long-term roadside data [2] will test cross-day transfer, time-of-day conditioning and rare-event statistics; change detection for stale memory, false alarms and spoofed objects, and a hybrid camera–fleet memory complete the safety case.

8 Conclusion

Long-term memory of a place, learned by fixed infrastructure cameras, predicts what happens there later and, used as a planning prior, lets an occlusion-aware robotaxi be safer at the same speed—by about a quarter on a real Washington, DC block, more where pedestrian activity is concentrated. The same experiments show how it fails and how to bound the damage: memory may only add caution unless it is known to be fresh, and live information must expire. Autonomous driving need not be only a problem of a smarter car; part of the intelligence can live in the city, as long as the car treats it as a prior with an expiry date.

Data and code availability

TGSIM Foggy Bottom is public (U.S. DOT, CC0) [25]. The repository contains the full pipeline (data download, cleaning, memory, evaluation, simulator, fleet comparison), 30 unit tests and all results; Parts 1–3 reproduce in about twenty minutes on a laptop. [Repository URL to be added.]

Acknowledgements

[To be completed: supervisor, and a statement on the use of AI tools in line with HSE University regulations.]

References

- [1] Haibao Yu et al. DAIR-V2X: A large-scale dataset for vehicle-infrastructure cooperative 3D object detection. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2022.

- [2] Haibao Yu et al. V2X-Seq: A large-scale sequential dataset for vehicle-infrastructure cooperative perception and forecasting. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2023.
- [3] Yue Hu, Shaoheng Fang, Zixing Lei, Yiqi Zhong, and Siheng Chen. Where2comm: Communication-efficient collaborative perception via spatial confidence maps. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2022.
- [4] Haibao Yu et al. End-to-end autonomous driving through V2X cooperation. In *Proceedings of the AAAI Conference on Artificial Intelligence*, 2025.
- [5] ETSI. Intelligent transport systems (ITS); vehicular communications; basic set of applications; collective perception service. Technical Report TS 103 324 V2.1.1, European Telecommunications Standards Institute, 2023.
- [6] ETSI. Intelligent transport systems (ITS); vehicular communications; basic set of applications; local dynamic map (LDM). Technical Report EN 302 895 V1.1.1, European Telecommunications Standards Institute, 2014.
- [7] Swaminathan Gopalswamy and Sivakumar Rathinam. Infrastructure enabled autonomy: A distributed intelligence architecture for autonomous vehicles. In *IEEE Intelligent Vehicles Symposium (IV)*, 2018.
- [8] Caixin Global. ‘vehicle-road-cloud integration’ pilot program to include 20 cities. <https://www.caixinglobal.com/2024-07-04/vehicle-road-cloud-integration-pilot-program-to-include-20-cities-102212883.html>, 2024.
- [9] Jiaxun Zhang et al. Cooperative safety intelligence in V2X-enabled transportation: A survey. arXiv:2512.00490, 2025.
- [10] Siyuan Meng and Chengbo Ai. Infrastructure-centric world models: Bridging temporal depth and spatial breadth for roadside perception. arXiv:2604.17651, 2026.
- [11] Mobileye. REM — road experience management. <https://www.mobileye.com/technology/rem/>, 2026. Accessed 29 September 2026.
- [12] Ross Greer, Lulua Rakla, Samveed Desai, Afnan Alofi, Akshay Gopalkrishnan, and Mohan Trivedi. CHAMP: Crowdsourced, history-based advisory of mapped pedestrians for safer driver assistance systems. arXiv:2301.05842, 2023.
- [13] Yurong You, Katie Z. Luo, Xiangyu Chen, Junan Chen, Wei-Lun Chao, Wen Sun, Bharath Hariharan, Mark Campbell, and Kilian Q. Weinberger. Hindsight is 20/20: Leveraging past traversals to aid 3D perception. In *International Conference on Learning Representations (ICLR)*, 2022.
- [14] Xuan Xiong et al. Neural map prior for autonomous driving. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2023.
- [15] Tomasz Piotr Kucner et al. Survey of maps of dynamics for mobile robots. *The International Journal of Robotics Research*, 2023.
- [16] Yufei Zhu et al. CLiFF-LHMP: Using spatial dynamics patterns for long-term human motion prediction. In *IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, 2023.

- [17] Yufei Zhu, Andrey Rudenko, Tomasz P. Kucner, Achim J. Lilienthal, and Martin Magnusson. Long-term human motion prediction using spatio-temporal maps of dynamics. arXiv:2510.03031, 2025.
- [18] Chi Zhang, Florian Steinhauser, Gereon Hinz, and Alois Knoll. Occlusion-aware planning for autonomous driving with vehicle-to-everything communication. *IEEE Transactions on Intelligent Vehicles*, 9(1):1229–1242, 2024.
- [19] Mert Koc, Ekim Yurtsever, Keith Redmill, and Ümit Özgüner. Pedestrian emergence estimation and occlusion-aware risk assessment for urban autonomous driving. In *IEEE International Intelligent Transportation Systems Conference (ITSC)*, 2021.
- [20] Derq. Derq INSIGHT: Safety insights and traffic analytics. <https://en.derq.com/safety-and-traffic-insights>, 2026. Accessed 29 September 2026.
- [21] Bo Shang and Yiqiao Li. Roadside LiDAR for cooperative safety auditing at urban intersections: Toward auditable V2X infrastructure intelligence. arXiv:2604.10419, 2026.
- [22] Xintao Yan et al. Learning naturalistic driving environment with statistical realism. *Nature Communications*, 14, 2023.
- [23] Yves-Alexandre de Montjoye, César A. Hidalgo, Michel Verleysen, and Vincent D. Blondel. Unique in the crowd: The privacy bounds of human mobility. *Scientific Reports*, 3:1376, 2013.
- [24] John C. Hayward. Near-miss determination through use of a scale of danger. *Highway Research Record*, 384, 1972.
- [25] Federal Highway Administration. Third generation simulation data (TGSIM) foggy bottom trajectories. U.S. Department of Transportation, <https://data.transportation.gov/d/brzy-6zfh>, 2024.
- [26] Julian Bock, Robert Krajewski, Tobias Moers, Steffen Runde, Lennart Vater, and Lutz Eckstein. The inD dataset: A drone dataset of naturalistic road user trajectories at German intersections. In *IEEE Intelligent Vehicles Symposium (IV)*, 2020.